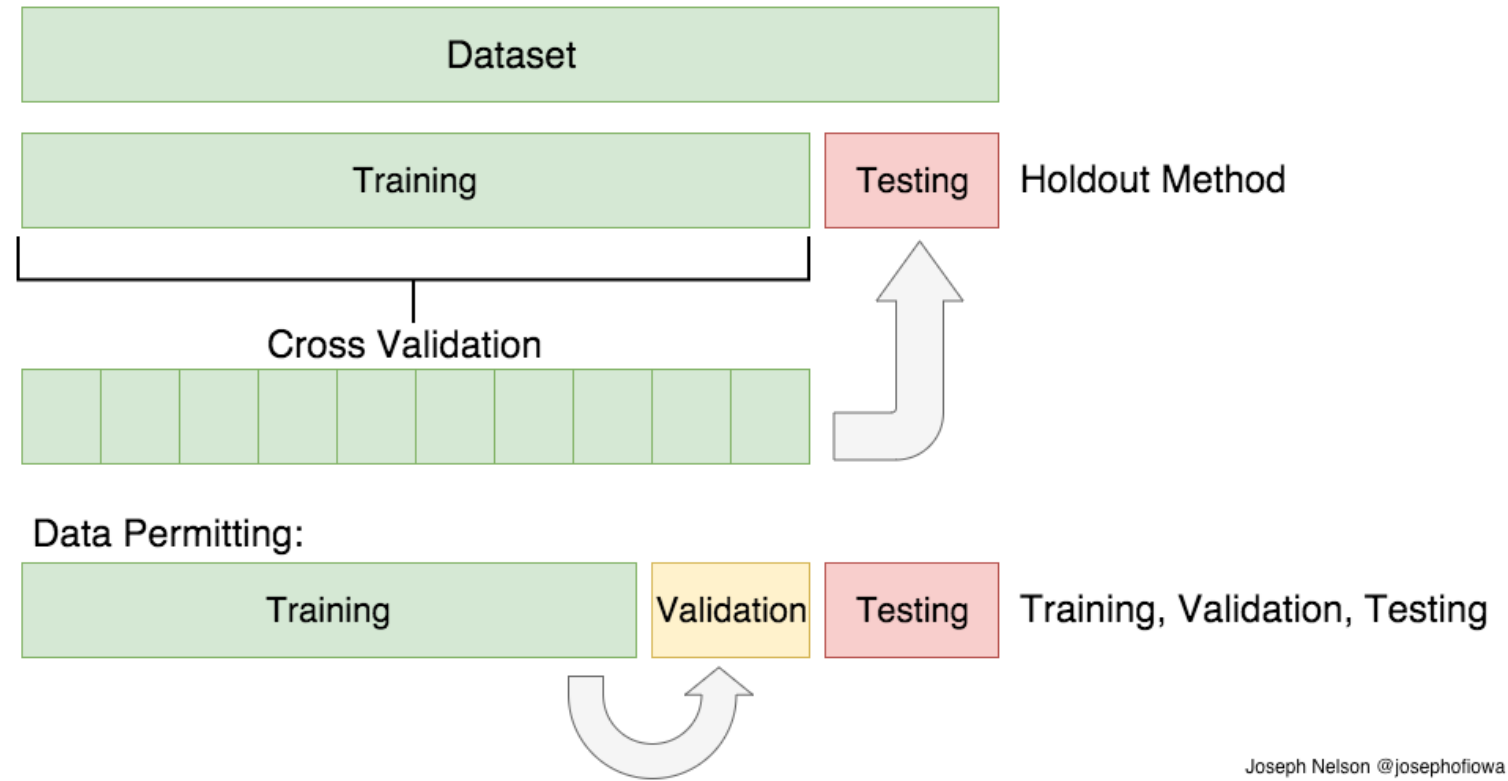



Cross-Validation

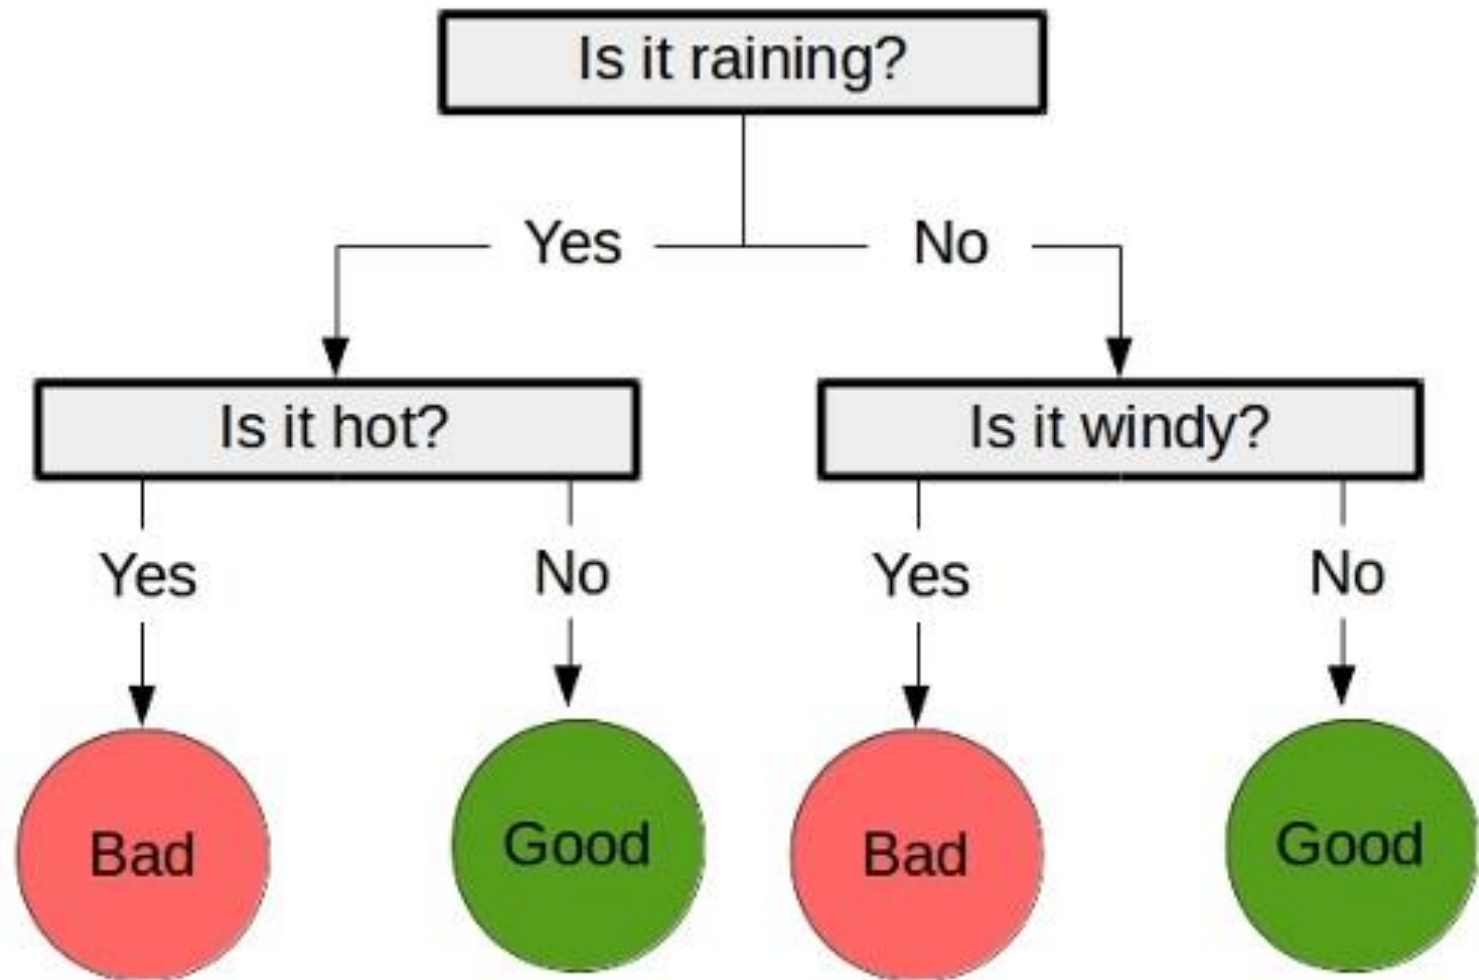
- Split your data
- Control overfitting
- Control underfitting
- Hyperparameters tuning



Linear Regression



Decision Tree

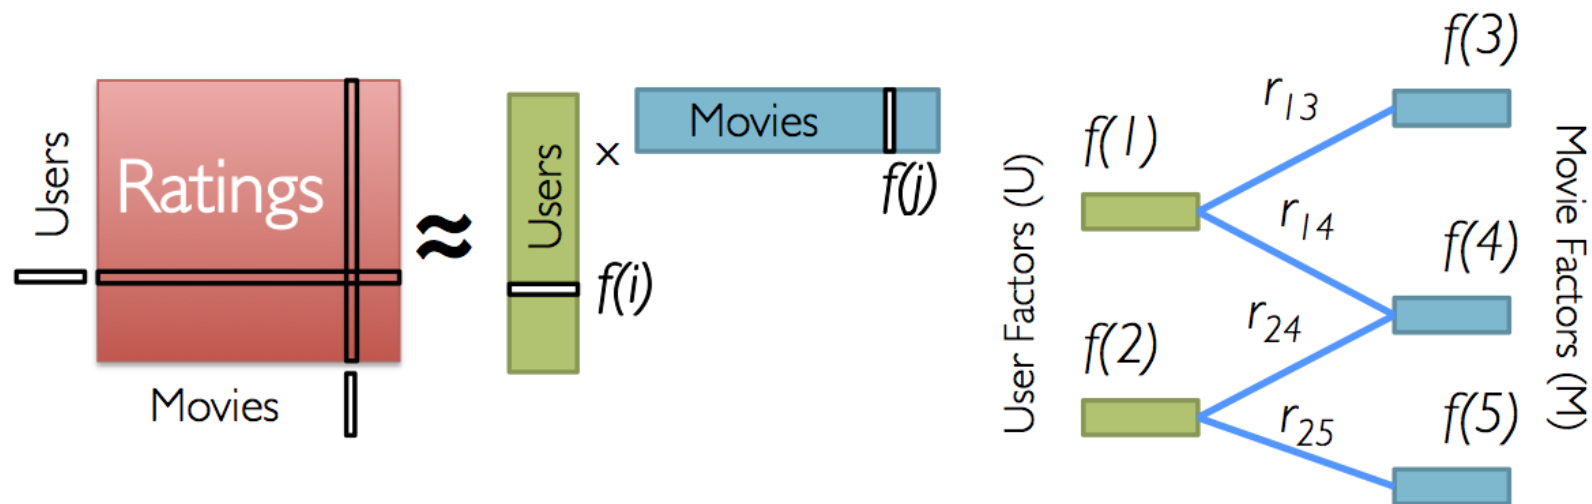


Some Results

Model	RMSE	R2
Linear Regression	0.455340	0.363434
Linear Regression with CV	0.446816	0.387042
Decision Tree with CV	0.372874	0.573129
Random forest	0.366229	0.588207
Gradient-Boosting DT	0.355057	0.612949

Recommendation
system

Low-Rank Matrix Factorization:



Iterate:

$$f[i] = \arg \min_{w \in \mathbb{R}^d} \sum_{j \in \text{Nbrs}(i)} (r_{ij} - w^T f[j])^2 + \lambda ||w||_2^2$$

What else?

- <https://docs.databricks.com/index.html>



Keras



PYTORCH

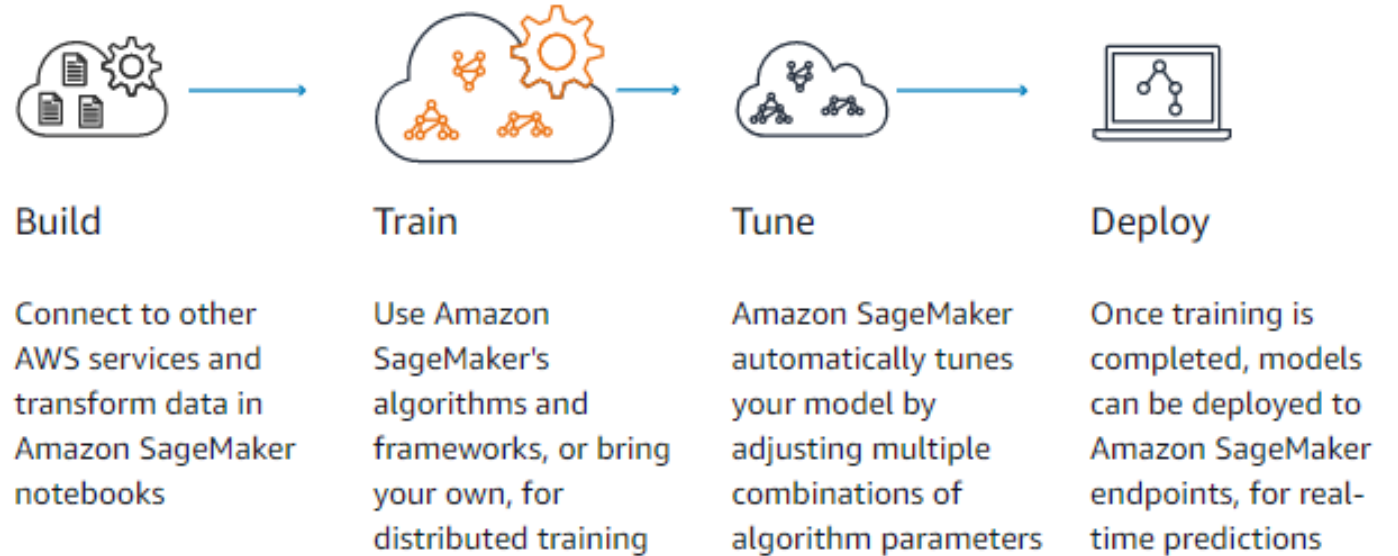
Deep Learning with PyTorch



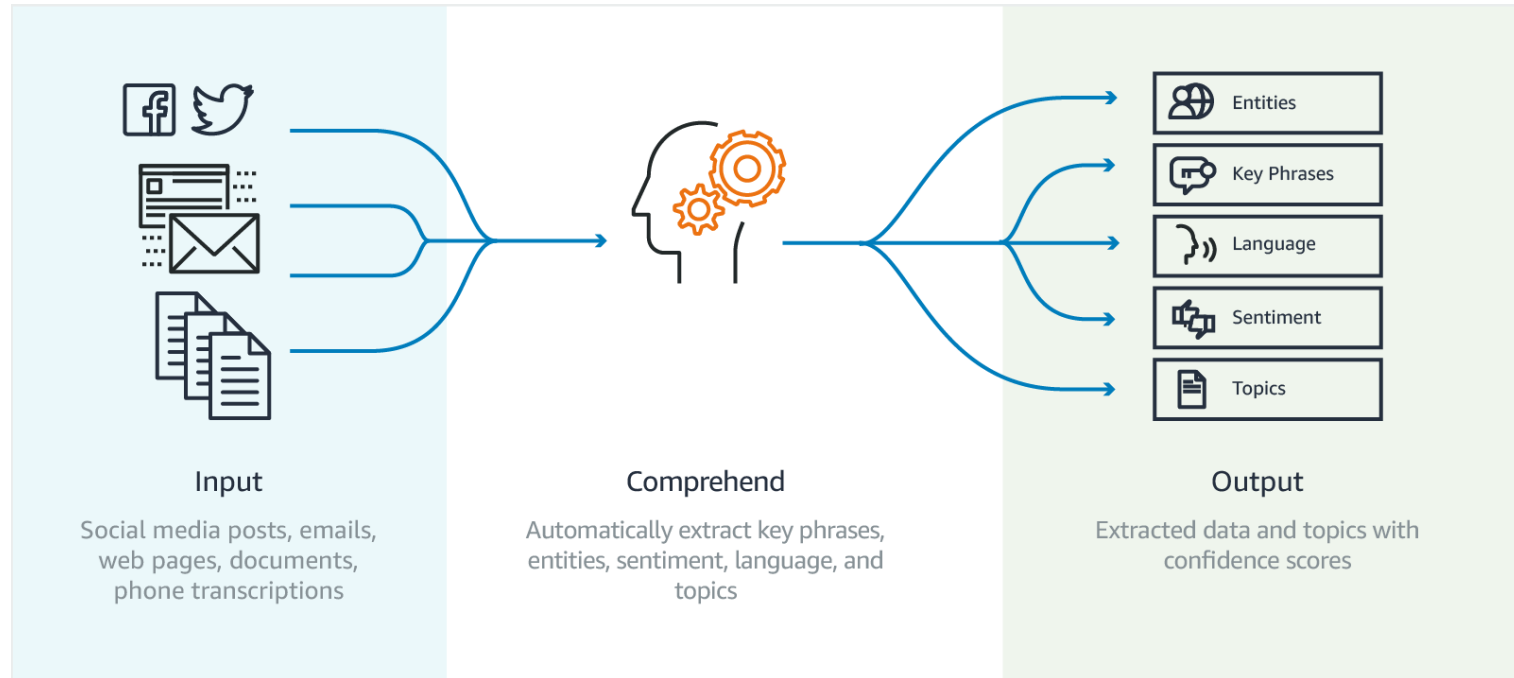


Amazon Machine Learning


Amazon SageMaker




Amazon Comprehend

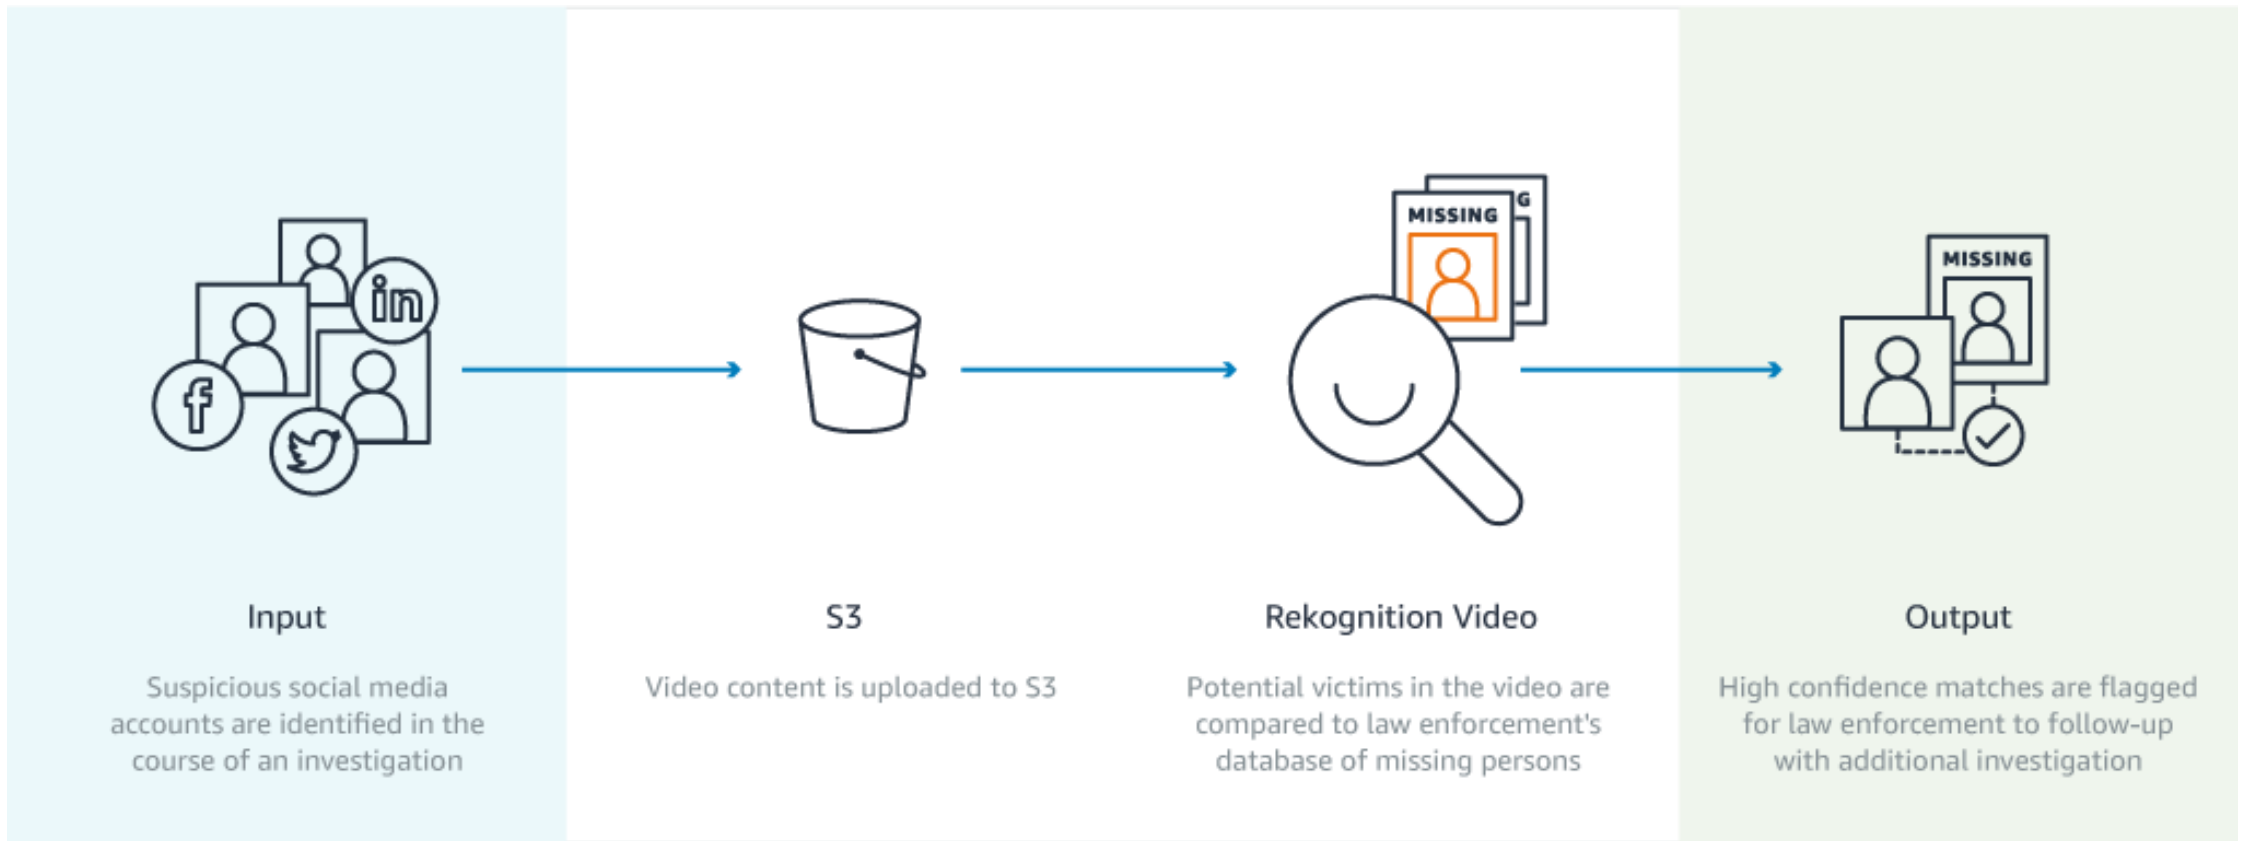


Amazon Lex





Amazon Polly



Amazon Rekognition



Machine learning

Amazon Transcribe

Automatic Speech Recognition

Amazon Transcribe provides high-quality and affordable speech-to-text transcription for a wide range of use cases.

Machine Translation Use Cases

Enable multilingual sentiment analysis of social media content

With Amazon Translate, you are not restricted by language barrier. Understand the social sentiment of your brand, product, or service while monitoring online conversations in different languages. Simply translate the text to English before using a natural language processing (NLP) application like Amazon Comprehend to analyze textual content in a multitude of languages.

Provide on-demand translation of user-generated content

It's very difficult for human translation teams to keep up with dynamic or real-time content. With Amazon Translate, you can easily translate massive volumes of user-generated content in real-time. Websites and applications can automatically make content such as feed stories, profile descriptions, and comments, available in the user's preferred language with a click of a "translate" button.


Add real-time translation for communications applications

Amazon Translate can provide automatic translation to enable cross-lingual communications between users for your applications. By adding real-time translation to chat, email, helpdesk, and ticketing applications, an English-speaking agent or employee can communicate with customers across multiple languages.

The image shows the Kaggle logo, which consists of the word "kaggle" in a blue, lowercase, sans-serif font. A small "TM" trademark symbol is located at the top right of the letter "e". To the left of the text is a thin, vertical blue line.

kaggle™

Platforms and Resources You Have Used to Continue Learning Data Science Skills




Competitions Platforms

- Kaggle
- DrivenData
- CrowdAnalitix
- CodaLab
- DataScienceChallenge.net
- DataScience.net
- Single-competition sites (like KDD, VizDooM)



Why to participate?

- Great opportunity to **learn** and networking
- Interesting non-trivial tasks and state-of-the-art approaches
- A way to become famous inside data science community
- A way to earn some money
- Wrap up competition results as:
 - LinkedIn project
 - GitHub repository
 - Medium blog





House Prices: Advanced Regression Techniques

Predict sales prices and practice feature engineering, RFs, and gradient boosting

5,058 teams · 2 years to go

[Overview](#)[Data](#)[Kernels](#)[Discussion](#)[Leaderboard](#)[Rules](#)[Team](#)[My Submissions](#)[Submit Predictions](#)

Overview

Description

[Evaluation](#)[Frequently Asked Questions](#)[Tutorials](#)

Start here if...

You have some experience with R or Python and machine learning basics. This is a perfect competition for data science students who have completed an online course in machine learning and are looking to expand their skill set before trying a featured competition.

Competition Description



